

# Semaphore for Synchronization

- *Race condition*: A situation that several tasks access and manipulate the same data concurrently and the outcome of the execution depends on *the particular order in which the access take place*.
- Example:
  - Suppose that the value of the variable **counter** = 5.
  - Process 1 and process 2 execute the statements “**counter++**” and “**counter--**” concurrently.
  - Following the execution of these two statements, the value of the variable **counter** may be 4, 5, or 6!

“**counter++**” is implemented as

```
register1 = counter
register1 = register1 + 1
counter = register1
```

“**counter--**” is implemented as

```
register2 = counter
register2 = register2 - 1
counter = register2
```

- The concurrent execution of “counter++” and “counter--” is equivalent to a sequential execution where the low-level statements are interleaved in some arbitrary order.

Process 1	register <sub>1</sub>	Process 2	register <sub>2</sub>	counter
register <sub>1</sub> = counter	5	---	---	5
register <sub>1</sub> = register <sub>1</sub> + 1	6	---	---	5
---	---	register <sub>2</sub> = counter	5	5
---	---	register <sub>2</sub> = register <sub>2</sub> - 1	4	5
counter = register <sub>1</sub>	6	---	---	6
---	---	counter = register <sub>2</sub>	4	4

# Shared memory synchronization

- There are two essential needs for synchronization between multiple processes executing on shared memory
  - **Establishing an order between two events**
    - E.g. in the server and client case, we want to make sure the server finishes writing before the client reads
  - **Mutually exclusive access to a certain resource**
    - Such as a data structure, a file, etc
    - E.g. Two people deposit to the same account “deposit +=100”. We want to make sure that the increment happens one at a time. Why? (Let us look draw a time line showing possible interleaving of events)

- A semaphore can be used for both purposes
- An ordinary while loop (busy wait loop) is not safe for ensuring mutual exclusion
  - Two processes may both think they have successfully set the lock and, so, have the exclusive access
    - Again, we can draw a time line showing possible interleaving of events that may lead to failed mutual exclusion
  - A semaphore is guaranteed to be able to have the correct view of the locking status

# The concept of semaphores

- Semaphores may be *binary* (0/1), or *counting*
- Every semaphore variable,  $s$ , It is initialized to some positive value
  - 1 for a binary semaphore
  - $N > 1$  for a counting semaphore

# Binary semaphores

- A binary semaphore ,  $s$  , is used for mutual exclusion and wake up sync
  - 1 == unlocked
  - 0 == locked
- $s$ , is associated with two operations:
- $P(s)$ 
  - Tests  $s$ ; if positive, resets  $s$  to 0 and proceed; otherwise, put the executing process to the back of a waiting queue for  $s$
- $V(s)$ 
  - Set  $s$  to 1 and wake up a process in the waiting queue for  $s$

# Counting semaphores

- A counting semaphore,  $s$ , is used for producer/consumer sync
  - $n$  == the count of available resources
  - $0$  == no resource (locking consumers out)
- $s$ , is associated with two operations:
- $P(s)$ 
  - Tests  $s$ , if positive, decrements  $s$  and proceed
  - otherwise, put the executing process to the back of a waiting queue for  $s$
- $V(s)$ 
  - Increments  $s$ ; *wakes up a process, if any, in the waiting queue for  $s$*



Process 1	register <sub>1</sub>	Process 2	register <sub>2</sub>	counter
<code>sem_wait(&amp;semA);</code>		---		5
---		<code>sem_wait(&amp;semA);</code>		5
<code>register<sub>1</sub> = counter</code>	5	<code>/* blocked */</code>	---	5
<code>register<sub>1</sub> = register<sub>1</sub> + 1</code>	6	<code>/* blocked */</code>	---	5
<code>counter = register<sub>1</sub></code>	6	<code>/* blocked */</code>		6
<code>sem_post(&amp;semA);</code>		<code>/* blocked */</code>		6
	---	<code>register<sub>2</sub> = counter</code>	6	6
	---	<code>register<sub>2</sub> = register<sub>2</sub> - 1</code>	5	6
	---	<code>counter = register<sub>2</sub></code>	5	5
		<code>sem_post(&amp;semA);</code>		

# Critical Sections

- We like to think of locking a concurrent data structure
- In current practice, however, locks (incl. binary semaphores) are typically used to lock a segment of program statements (or instructions)
- Such a program segment is called a *critical section*
  - *A critical section is a program segment that may modify shared data structures*
  - *It should be executed by one process at any given time*

- With a binary semaphore
  - If multiple processes are locked out of a critical section
    - As soon as the critical section is unlocked, only one process is allowed in
    - The other processes remain locked out
- Implementation of semaphores is fair to processes
  - A first-come-first-serve queue

# Unix Semaphores

- There are actually at least two implementations
- UNIX System V has an old implementation
  - Analogous to shared memory system calls
  - Calls to `semget()`, `semat()`, `semctl()`, etc
  - Not as easy to use as POSIX implementation
- We will use POSIX implementation in this course

# POSIX semaphore system calls

- `#include <semaphore.h>`
- POSIX semaphores come in two forms: named semaphores and unnamed semaphores.

# Using unnamed semaphores

- Unnamed semaphores are also called memory-based semaphores
  - Named semaphores are “file-based”
- An unnamed semaphore does not have a name.
  - It is placed in a region of memory that is shared between multiple threads (a thread-shared semaphore) or **processes** (a process-shared semaphore).
- A process-shared semaphore must be placed in a shared memory region

# System calls

- Before being used, an unnamed semaphore must be initialized using [sem\\_init\(3\)](#). It can then be operated on using [sem\\_post\(3\)](#) and [sem\\_wait\(3\)](#).
- When the semaphore is no longer required, and before the memory in which it is located is deallocated, the semaphore should be destroyed using [sem\\_destroy\(3\)](#).
- Compile using -lrt

Recall that shared memory segments must be removed before program exits

- “An unnamed semaphore should be destroyed with `sem_destroy()` before the memory in which it is located is deallocated.”
- “ Failure to do this can result in resource leaks on some implementations.”



```
int sem_init(sem_t *sem, int pshared, unsigned  
int value);
```

- #include <[semaphore.h](#)>
- **sem\_init()** initializes the unnamed semaphore at the address pointed to by *sem*. The *value* argument specifies the initial value for the semaphore.
- If *pshared* has the value 0, then the semaphore is shared between the threads of a process
- If *pshared* is nonzero, then the semaphore is shared between processes, and should be located in a region of shared memory

```
int sem_wait(sem_t *sem);
```

- **sem\_wait()** decrements (locks) the semaphore pointed to by *sem*.
- If the semaphore's value is greater than zero, then the decrement proceeds, and the function returns, immediately.
- If the semaphore currently has the value zero, then the call blocks until either it becomes possible to perform the decrement (i.e., the semaphore value rises above zero), or a signal handler interrupts the call.

```
int sem_post(sem_t *sem);
```

- **sem\_post()** increments (unlocks) the semaphore pointed to by *sem*.
- If the semaphore's value consequently becomes greater than zero, then another process or thread blocked in a **sem\_wait**(3) call will be woken up

**int sem\_destroy(sem\_t \*sem);**

- **Destroys** the unnamed semaphore at the address pointed to by *sem*. Only a semaphore that has been initialized by [sem\\_init\(3\)](#) should be destroyed using **sem\_destroy()**.
- Destroying a semaphore that other processes or threads are currently blocked on (in [sem\\_wait\(3\)](#)) produces undefined behavior.
- Using a semaphore that has been destroyed produces undefined results, until the semaphore has been reinitialized using [sem\\_init\(3\)](#).

# Examples

- We first look at a bad example in which the unnamed semaphore is not placed in the shared memory (test1.c)

```
// compile with -lrt  
#include <semaphore.h>
```

```
#include <stdio.h>  
#include <errno.h>  
#include <stdlib.h>  
#include <unistd.h>  
#include <sys/types.h>  
#include <sys/shm.h>  
#include <sys/wait.h>  
#define SHMSIZE 1024  
int main(int argc, char **argv)  
{  
    int i,nloop=10,*ptr;  
    sem_t mutex;
```

```
    int shmid1;  
    int *shm1, *s;  
    if ((shmid1 = shmget(IPC_PRIVATE,  
SHMSIZE, 0666)) < 0) {  
        perror("shmget");  
        exit(1);  
    }  
    if ((shm1 = shmat(shmid1, NULL, 0)) == (int  
*) -1) {  
        perror("shmat");  
        exit(1);  
    }  
    *shm1 = 0;  
    ptr = shm1;
```

- In this example, the semaphore is not placed in the shared memory.
- Therefore, it is ineffective for mutual exclusion synchronization

```

/* create, initialize semaphore */
if( sem_init(&mutex,1,1) < 0)
{
    perror("semaphore
initilization");
    exit(0);
}
if (fork() == 0) { /* child process*/
    sem_wait(&mutex);
    for (i = 0; i < nloop; i++) {
        printf("child: %d\n", (*ptr)++);
        sleep(5);
    }
    sem_post(&mutex);
    exit(0);
}

```

```

/* back to parent process */
sem_wait(&mutex);
for (i = 0; i < nloop; i++) {
    printf("parent: %d\n",
(*ptr)++);
    sleep(5);
}
sem_post(&mutex);
wait(int *) 0);
shmctl(shmid1, IPC_RMID,
(struct shmid_ds *) 0);
exit(0);
}

```

- The mutex is supposed to ensure that each process prints its entire data w/o mixing with the other process' data
- But it fails to do so

- Next, we look at an even worse example:
  - We want to let parent process prints its entire data first
  - So we let child process wait for the process to give it a go-ahead
  - Initialize the mutex variable to 0 and wait for the parent process to change it to 1.
- But we didn't put the mutex variable in the shared memory
- The child process never wakes up!
- We need to manually kill the child process and free the shared memory



```

#include ..... // stuck.c
int main(int argc, char **argv)
{
    int i,nloop=10,*ptr;
    sem_t mutex;
    .....
    if( sem_init(&mutex,1,1) <
0) */
    if( sem_init(&mutex,1,0) <
0) { .....
    }
    if (fork() == 0) { /* child
        process*/ sem_wait(&mutex);
        for (i = 0; i < nloop; i++)
            printf("child: %d\n", (*ptr)++);
        exit(0)
    }
}

```

```

/* back to parent
process */
for (i = 0; i < nloop;
i++)    printf("parent:
%d\n", (*ptr)++);

    sem_post(&mutex);
    exit(0);

```

- Finally, we will correct the errors by placing the semaphore in the shared memory
- We also need to remember to destroy the unnamed semaphore before removing the shared memory segment.
- Be careful with the timing for destroying the semaphore
  - Make sure there should not be waiting processes

```
// nonstuck.c
sem_t *p_mutex;

.....
if ((shmid2 = shmget(IPC_PRIVATE, SHMSIZE, 0666)) < 0)
    { perror("shmget");
      exit(1);
    }
p_mutex = (sem_t *) shmat(shmid2, NULL,
0); if (p_mutex == (sem_t *) -1) {
    perror("mutex shmat fails
"); exit(1);
}
if( sem_init(p_mutex,1,0) < 0)
{
    perror("semaphore initialization");
    exit(1);
}
if (fork() == 0) { /* child process*/
    sem_wait(p_mutex);
```

// cont'd on next page

```
// nonstuck.c cont'd
if (fork() == 0) { /* child process*/
    sem_wait(p_mutex);
    for (i = 0; i < nloop; i++)
        printf("child: %d\n", (*ptr)++);
    sem_destroy(p_mutex);
    shmctl(shmid2, IPC_RMID, (struct shmid_ds *)
0); shmctl(shmid1, IPC_RMID, (struct shmid_ds
*) 0); exit(0);
}
/* back to parent process */ for
(i = 0; i < nloop; i++)
    printf("parent: %d\n", (*ptr)++);
sem_post(p_mutex);
exit(0);
```

- We can make a similar change to test1.c
- We will see that now each process will print its entire data without interleaving with other processes
- Which process writes first will be unknown in advance